


27. System Infrastructure Overview

The Cloud Products Engineering (CPE) division maintains the core infrastructure supporting all Aetheris AI platforms. Two flagship internal initiatives — Project Falcon and Project AetherOS — share a common request path that begins at the AetherFlow API gateway and terminates at the primary inference and database clusters described below.

27.1 Request Path

Incoming client requests are authenticated at the API Gateway, distributed by the Load Balancer, and routed to either the inference cluster (for model serving) or the database server (for structured lookups). Model weights and large artifacts are served from Object Storage on demand.

Aetheris Core Infrastructure — Project Falcon / AetherOS




27.2 Key Server Identifiers

TITAN-ML-08 — primary inference cluster serving Project Falcon transformer workloads.
ZEPHYR-DB-01 — primary structured-data database server underpinning Project AetherOS.

28. Building Floor & Security Zone Map

The Seattle Aetheris Tower spans floors 20 through 35. Access is segmented into color-coded security zones to enforce the Visitor Policy (Section 16) and IT Security Policy (Section 10). The diagram below summarizes floor assignments referenced elsewhere in this handbook.




28.1 Zone Summary

Green Zone (Floor 20): Visitor check-in and temporary badge issuance.
Standard Zone (Floors 22-30): Bookable meeting rooms and badge-release printing stations.
Red Zone (Floors 34-35): AI Research Lab server rooms; visitor access is prohibited at all times.

29. Support Escalation Flow Diagram

This diagram illustrates the escalation path referenced in the Customer Support Process (Section 25). Tickets are triaged into severity tiers immediately after ingestion, each governed by a distinct SLA. Unresolved L1 tickets are automatically escalated after 3 hours.

Global Customer Success — Support Escalation Flow



29.1 SLA Reference Table (Recap)

- L1 (Critical): acknowledged within 15 minutes; resolved within 4 hours; escalates to the VP of Engineering, Sarah Jenkins, if unresolved after 3 hours.
- L2 (Major): acknowledged within 2 hours; resolved within 24 hours.
- L3 (Minor): handled via the standard support queue with no fixed escalation timer.